A free archive.
A negative first test.

We obtained real aggregate transport data and tested it against the requested military cases. This particular measure does not yet justify a newsroom warning service.

What was actually tested

Pentagon Metro station entries from 7 p.m. to midnight, expressed as a share of afternoon-plus-evening entries. Each date was compared with earlier matching weekdays and three nearby stations. The question was whether rail departures shifted unusually late before an operation. These are passenger passages, **not employees, unique people or building occupancy**.

Military case	Pre-event test	Friday evening entries
Iran: June 2025	No flag; 2 of 3 evenings eligible	245 vs. 298 usual*
Venezuela: January 2026	Unscorable: holiday / low volume	133 vs. 250 usual*
Iran: February 2026	No flag; 3 of 3 evenings eligible	314 vs. 250 usual*

* Prior comparable-Friday median, rounded here. Friday counts are descriptive; the screening rule used a three-evening window and comparison stations. Venezuela is not counted as a failed prediction.

The background test matters

The threshold was fixed using 2024. Applied unchanged from January 2025 through March 2026, it flagged **23 of 259 eligible dates**, grouped into 18 alert episodes. That is 2.66 flags per 30 scored dates, above the chosen workload target of two. The full interval contained 455 dates; 196 were excluded or lacked a valid score.

These are background alerts, not proven false military predictions: other events were not exhaustively classified. With two scored military cases and one unscorable case, no general accuracy estimate is warranted.

Decision: retain this as a reproducible benchmark; do not deploy this Metro measure alone. A negative result for rail departures does not rule out other aggregate activity signals.

What the pilot makes possible

A real dataset, at zero acquisition cost

The kit contains 12 original exports and a cleaned panel of **14,688 records**: four stations, three time selections and 1,224 dates from January 1, 2023 to May 8, 2026. There were no missing date rows or blank count cells. The afternoon and evening subtotals passed consistency checks against daily totals. Source accuracy remains a separate question.

The comparisons are Pentagon City, Rosslyn and Crystal City. Downloads, cleaning, scoring and a synthetic sensitivity check are reproducible using ordinary Python, without paid packages, an account or an API key. The archive is public; no individual movements or device identifiers were collected.

The binding constraint is publication delay

WMATA documents a **seven-day lag** for this archive [1]. Even a strong retrospective pattern would not demonstrate a timely public warning. The records were downloaded now; they are not a versioned archive of exactly what was published before each event.

Time and ordinary activity can create convincing mistakes

Metro service dates run from 4 a.m. to 3:59 a.m. the following day [1]. A Friday daily total can therefore include early Saturday. The test used separately exported afternoon and evening blocks. June 2025 impacts were on June 21 in Washington, although commonly dated June 22 in Iran. Venezuela's January 2-3 operational timeline still needs tighter verification; a stricter cutoff also left that window unscorable.

Documented holidays, the 2025 federal shutdown, selected severe-weather days and February 2026 construction closures were excluded where relevant. Crystal City was closed on the February Iran operation date [2]; unadjusted post-event comparisons could be misleading. The disruption calendar remains incomplete. The 1,000-passage scoring floor is an analyst choice for broad traffic, not a measure of staff or anonymity.

What should happen next

Use the included code for independent replication. A fuller review needs a complete disruption calendar and a military-event registry chosen independently of the signal. Do not tune the rule until these familiar cases look predictive. Before any live pilot, verify a different aggregate channel's historical coverage, actual publication times, access and cost. This audit has not established a feed meeting all those requirements.

Sources and reproducibility

[1] WMATA: How to Use and Interpret Metro Ridership Data (July 30, 2026).
[2] WMATA: February 2026 Ridership Snapshot.
Full source URLs, event references, raw data, exact numerical results, exclusions and code are in the accompanying pilot kit. Figures in this report are our exploratory calculations.